

# Shoot for the Stats, Alm for the Moon:

Developing a pixel Almbot


Alessio Barboni - Francesco Redaelli

# Project aim

Developing an *AI agent* able to  
play a *FPS Aim Trainer game*


The agent would exploit only visual information (i.e., screen pixels color values) in order to properly move the mouse pointer  
live, in a *human-mimicking fashion*

# The game level



- Goal:  
Chase the target
- Duration:  
60 seconds
- Score:  
# target hits

# Interlude: mouse position on the screen



*In our project:*

$$mp = (X, Y)$$

$$X \in [0, 1920]$$

$$Y \in [0, 1080]$$

# Data Collection

We played the target level multiple times and collected a data point every 1/20th of a second

## Data Point (time t)



*Screenshot<sub>t-1</sub> (1920x1080)*



- Mouse Movement on the X axis:  $X_t - X_{t-1}$ 
  - $\text{Int}(\rho x) \in [-1920, 1920]$
- Mouse Movement on the Y axis:  $Y_t - Y_{t-1}$ 
  - $\text{Int}(\rho y) \in [-1080, 1080]$

# Data Collection “Hack”

Although we aimed at training a model that would reproduce our *far-from-being-perfect playstyle*, we nonetheless laid down some *data collecting rules* in an attempt to improve the quality of the input data:

- Do not miss!
  - Only *100%* shooting accuracy levels data saved
- Move!
  - No standing still data point - *(0,0)* movement vector - allowed

# Data preprocessing

# Data preprocessing - Masking

BGR Boundaries




Lower = [0,150,150]

Upper = [255,255,255]



*Original Screenshot*




*Masked Screenshot*




# Data preprocessing - From BRG to Grayscale

$$Y = 0.299 R + 0.587 G + 0.114 B$$




*Masked Screenshot*




*Grayscale Screenshot*

# Data preprocessing - Removing clock & time

Blackened the corresponding pixel (set their value to 0)



*Grayscale Screenshot*




*Preprocessed Screenshot*

Mouse movement:  
regression model

# Data preprocessing - Regression: resizing

New size: (6x6) - Grid search approach


Algorithm: INTER\_AREA (resampling using pixel area relation)




*Preprocessed Screenshot (1920x1080)*

*Resized Screenshot (6x6)*


# Data augmentation - Regression




Original  
( $X_{vec}$ ,  $Y_{vec}$ )




Flip around X-axis  
( $X_{vec}$ ,  $-Y_{vec}$ )




Flip around Y-axis  
( $-X_{vec}$ ,  $Y_{vec}$ )




Flip around both axis  
( $-X_{vec}$ ,  $-Y_{vec}$ )




Rotate 90°  
CounterClock  
( $Y_{vec}$ ,  $-X_{vec}$ )



Rotate 90°  
Clock  
( $-Y_{vec}$ ,  $X_{vec}$ )



Flip around X-axis &  
Rotate 90° CounterClock  
( $-Y_{vec}$ ,  $-X_{vec}$ )



Flip around X-axis  
& Rotate 90° Clock  
( $Y_{vec}$ ,  $X_{vec}$ )

# Data preprocessing - Regression: flattening & scaling

0	1	2	3	4	5
6	7	8	9	10	11
12	13	14	15	16	17
18	19	20	21	22	23
24	25	26	27	28	29
30	31	32	33	34	35



0	1	2	...	35
---	---	---	-----	----

$p_0$	$p_1$	$p_2$	...	$p_{35}$
-------	-------	-------	-----	----------

Pixels value  $\in [0,255]$



$\frac{p_0}{255}$	$\frac{p_1}{255}$	$\frac{p_2}{255}$	...	$\frac{p_{35}}{255}$
-------------------	-------------------	-------------------	-----	----------------------

Pixels value  $\in [0,1]$

# Target variables description

	x_mov	y_mov
count	20398.000	20398.000
mean	0.323	-0.208
std	59.260	50.298
min	-272.000	-170.000
25%	-25.000	-20.000
50%	0.000	0.000
75%	24.000	16.000
max	326.000	298.000

*Before data augmentation*

	x_mov	y_mov
count	163184.000	163184.000
mean	0.000	0.000
std	54.962	54.962
min	-326.000	-326.000
25%	-22.000	-22.000
50%	0.000	0.000
75%	22.000	22.000
max	326.000	326.000

*After data augmentation*

## Model Training: Train/Test Split

75% Train - 25% Test Size

# Mouse movement model: Random Forest regression

INPUT: 36 pixel values

OUTPUT: Two-dimensional vector (X\_mov, Y\_mov)

Multi-Output Strategy:

Fitting **one regressor per target** (sharing the *same parameters*, each one having as target a *different output component*: X\_mov and Y\_mov)

Score the model according to the **arithmetic average** of the individual regressor scores (different scoring strategies might be possible)

Random Forest Fitting: 5-Fold CV

Parameters:  $n_{estimators} = 100$

Scoring Metric:

$$RMSE = \sqrt{\sum_{i=1}^n \frac{(\hat{y}_i - y_i)^2}{n}}$$



# Mouse movement model: grid search approach

Input screenshot size	RMSE score
4,4	35.1
5,4	37.7
4,5	36.2
5,5	38.2
6,6	33.0
7,7	35.2
8,8	34.4
10,10	35.8


Best Model

*Input size = (6x6)*

*5-Fold CV RMSE = 33.0*



# Qualitative model assessment - Live Action!




# Interpretable Machine Learning (IML)

"Why Should I Trust You?"

An analysis based on the work and  
paper by *Marco Tulio Ribeiro et al.*  
(2016)

# Local Interpretable Model-agnostic Explanations (LIME)

Goal: *identify an interpretable model over the interpretable representation that is locally faithful to the classifier/regressor*



$$\xi(x) = \operatorname{argmin}_{g \in G} \mathcal{L}(f, g, \pi_x) + \Omega(g)$$

Explanation  
(intelligible)

Unfaithfulness  
(as opposed to  
local fidelity)

Complexity  
(as opposed to  
interpretability)


## LIME in action



RF Prediction:  $(-28.56, 1.54)$

Actual Response:  $(-84, 0)$

# LIME in action: X-axis



Feature Value


20	0.04
19	0.00
14	0.00
22	0.02
21	0.00
16	0.00
13	0.00
15	0.00
12	0.00
23	0.00
18	0.00
29	0.00
17	0.00

Intercept: -13.41


Score ( $R^2$ ): 0.296

LIME Prediction: -9.22

# LIME in action: X-axis

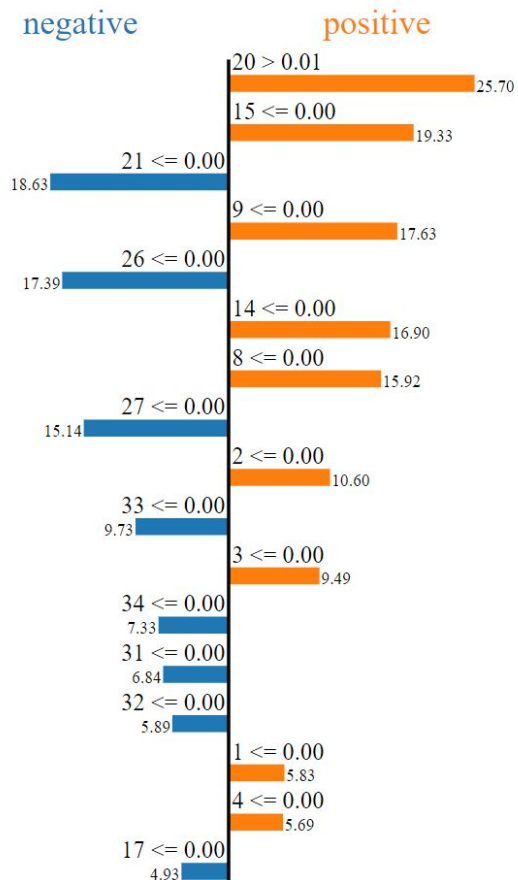



*Input Image*



*LIME Explanation*

# LIME in action: Y-axis



Feature Value

20	0.04
15	0.00
21	0.00
9	0.00
26	0.00
14	0.00
8	0.00
27	0.00
2	0.00
33	0.00
3	0.00
34	0.00
31	0.00


Intercept: -13.29

Score ( $R^2$ ): 0.269


LIME Prediction: 22.3



# LIME in action: Y-axis



*Input Image*



*LIME Explanation*

# Submodular Pick method for explaining models


Goal: pick a *diverse, representative set* of non-redundant explanations that represent how the model behaves globally

$$\text{Pick}(\mathcal{W}, I) = \underset{V, |V| \leq B}{\operatorname{argmax}} c(V, \mathcal{W}, I)$$

NP hard problem!



Greedy approach



*Explanation matrix  $W$ : represents the local importance of the interpretable components for each instance*

# I knew I should not have trusted you...

## Model issues:

- LIME explanations showed that the model at times failed in discerning whether a grey pixel in input corresponded to a **target** or to one of the **corners**
- Qualitative assessment highlighted some difficulties in reaching the actual target when the cursor was in its **neighborhood** right **above** or **below** it

## Our implemented preprocessing solution:

- *Double Masking + Binary Pixel Color Conversion*
- *Higher resolution input image*

# Data preprocessing - Double masking + BPCC



*Original Screenshot*

Binary Pixel Color Conversion

$$y = \begin{cases} 255 & \text{if } x > 0 \\ 0 & \text{if } x = 0 \end{cases}$$

BGR Boundaries

Upper = [0,0,252]  
Lower = [255,255,255]



Binary Pixel  
Color Conversion




BGR Boundaries


Upper = [240,100,0]  
Lower = [255,255,255]



Binary Pixel  
Color Conversion



*Target*




*Corners*

# Data preprocessing: Double Masking

*Formula:*  $Target + Corners * \frac{k}{255}$

$k = 100 \rightarrow$





*Final Screenshot*

# Data preprocessing - Regression: resizing

New size: (14x14) - Grid search approach

Algorithm: INTER\_AREA (resampling using pixel area relation)



*Preprocessed Screenshot (1920x1080)*

*Resized Screenshot (14x14)*

## Mouse movement model: grid search approach


Input screenshot size	RMSE score
10,10	30.1
12,12	29.3
14,14	28.8
16,16	29.0
20,20	29.4
24,24	29.7

Best Model

*Input size = (14x14)*

*5-Fold CV RMSE = 28.8*

# Qualitative model assessment - Live Action!






## LIME in action



RF Prediction:  $(-2.91, -19.96)$

Actual Response:  $(-10, -31)$

# LIME in action: X-axis



Feature Value


92	0.00
103	0.00
106	0.00
89	0.00
87	0.00
102	0.00
108	0.00
93	0.00
74	0.00
107	0.00
101	0.00
116	0.00
94	0.00

Intercept: -78.87


Score ( $R^2$ ): 0.598

LIME Prediction: -8.65

# LIME in action: X-axis



*Input Image*



*LIME Explanation*

# Conclusion

By analysing the explanations provided by the LIME IML algorithm and applying domain specific knowledge to modify the data preprocessing pipeline, we obtained:

- Better quantitative performance
  - *Lower RF RMSE*
- Better qualitative performance
  - *Smoother playstyle*
- Less of a “black box” and more trustworthy model
  - *Greater explainability and higher  $R^2$  score of LIME model*